



Stop Giving Models Open-Book Tests

Offensive Reasoning Index

How I got ORI ready to run benchmarks

Matthew Nickerson / SpecterOps

A model worth using

Matthew Nickerson


Adversary Simulation Consultant / SpecterOps

RED TEAM

LOCAL LLMs

BLOODHOUND

I want to know where the model is useful.
And where I still need to check its work.





My MCP needed a fair test

Which part of
the result belongs
to the model?

GOAD prompted the comparison.
I wanted randomized
evaluation data.

MODEL

Who is answering?

TOOLS

Which MCP or Direct?

GRAPH

What can it reach?

QUESTIONS

What are we asking?





Same seed. Same test.

67

ONE SEED

GRAPH


QUESTIONS



The first win was a very small question

EARLY EXAMPLE

Which user has an active session on
WS-IT-01.CORP.LOCAL?



Before reasoning, the data had to load

GENERATED JSON



SHARPHOUND ZIP



BLOODHOUND



UPLOAD FAIL

The file was not readable.

INGEST FAIL

Closer. Now fix fields, labels and edges.

A random seed does not stop the clock

LATER DETERMINISM BUG / FIELD NAMES


```
pwdlastset  
lastlogon  
lastlogontimestamp
```

Same seed
Different wall clock

Normalize volatile timestamps.
Fix ZIP-entry timestamps too.


Check the bytes, not just the seed.

From graph to grade




Two routes to the same graph

DIRECT CYPHER



Generate the query.
Execute it against the graph.
Grade the returned evidence.

MCP



Call → response → next step
Then submit the final answer.

The prompt is part of the experiment

HISTORICAL PHASE 2 / EXACT EXCERPT

Return ONLY the Cypher query –
no explanation, no markdown
fences, no commentary.

CURRENT MCP / EXACT EXCERPTS

Return exactly one JSON
object matching
submission_schema.

Do not invent graph evidence.

My grader missed correct answers

PROTOTYPE

EXECUTION

LATER REPAIRS

Exact text match



Query → results



Identity +
format

Missed some correct
answers

Compare returned
evidence

Normalize aliases; track
compliance



From local GPUs to hosted models

RTX 3080 -> M4 Mac -> Dual RTX 3090

Codex login → OpenAI models

OpenAI-compatible → NOUS / OpenRouter





When the easy questions stopped helping

One Cypher query can traverse multiple graph hops.

LOOKUP

Hand-checkable
questions

CONNECTED

Combine attack
techniques

COMPLEX

Sessions, pivots,
decoys, no-route tasks



From finding a path to explaining it

EARLIER / EXACT QUESTION


What is the attack path from
DJOHNSON@CORP.LOCAL to
DC01.CORP.LOCAL?
Show the full path.

Current State / Example Question


Find the viable multi-hop route.
Use multiple graph lookups.
Explain each mechanism.
Reach the Tier 0 condition.
Reject attractive dead ends.

Four hosts to Domain Admins

Initial host + three administrative pivots = four machines



Same path. Real graph.




A tool call in name only

PRINTED TEXT / ILLUSTRATION

```
{ "tool": "cypher_query" }
```

No invocation reached the server.



Record the actual tool events and the final answer.



The answer was there. Parsing hid it.

SAVED-ANSWER REPLAY / SAME ANSWERS, DIFFERENT EXTRACTION

RECORDED SCORE

3 / 50



OFFLINE EXTRACTION

34 / 50



Same saved answers. No new inference. No semantic repair.





Beyond the score

BENCHMARK TELEMETRY / METRICS + TRACES + RUN IDENTITY

OUTCOMES & COVERAGE

Scheduled / attempted / correct
Graded, invalid, errors
Timeouts and missing work

MODEL TRACES

Prompts and model turns
Exposed reasoning, answers
Generated queries

TOOL ACTIVITY

Call requests / responses
Tool names and arguments
Results and error flags

MCP RESOURCES

List and read interactions
Requests / responses / errors
URIs and returned content

USAGE & TIME

Reported input / output tokens
Task time and attempt history
Encode / decode rates*

RUN PROVENANCE

Model, provider and MCP route
Seed, repetitions, fingerprints
Task and transcript identities

* Provider-dependent detail. Resource responses $\neq$ successful loads.



Benchmarks on a budget

RUN LIMITS / REQUESTED SETTINGS, NOT FAILURE DIAGNOSES

OUTPUT TOKENS

8,192

Per model response

DIRECT

600s

Requested task deadline

MCP

1,200s

Requested task deadline

Selection target: roughly under \$5 / 1M output tokens

Smaller / Flash / affordable models; Anthropic not tested because of API cost.

■ A budget-focused comparison, not a frontier-model ranking.

A stopped run is not automatically a reasoning failure.





One run isn't the whole story

REPEAT SEED 67

50 fixed questions.
Three passes of the same test.
Compare scores and coverage.

FIXED BENCHMARK SCOPE

Seed 67 only.
No three-seed expansion.
No generalization claim.



ORI in action

Generating, Configuring, Running, Results

Early local scorecards

CORRECT / SCHEDULED – SEPARATE COHORTS, NOT A LEADERBOARD

[illegible]

- Different questions, configurations and graders.



Different tests, separate scorecards

SEPARATE SCORE PANELS / DIFFERENT CONTRACTS AND CORPORA

HISTORICAL V1 / ONE PASS

GPT-5.6 Sol 44/62



GPT-5.5 46/62



The grader had defects.

Historical shows how far it has come

HISTORICAL V29 / DIRECT

GPT-5.5 207/210



Daybreak Red 206/210



Daybreak Blue 205/210



GPT-5.6 Sol 204/210



42 tasks × 5 passes / model



The questions that never ran

LOCAL QWEN / ALL FIVE ROUTES / SEED 67

ROUTE	CORRECT	ATTEMPTED	RUN STATUS
Direct	29/50	50/50	completed
BloodHound MCP	34/50*	50/50	partial
MorDavid	2/50†	5/50	cutoff: 45 not run
Armadin	20/50*	50/50	partial
Steven external	35/50*	50/50	cleanup failed

 205/250 attempted. * failed/partial experiment † incomplete coverage



Steven: 35 correct + 9 incorrect + 5 invalid + 1 infrastructure; full task coverage, then cleanup failed [\[8\]](#).



Direct and MCP can disagree

SEED 67 / FULL-COVERAGE SCORES AND INFRASTRUCTURE CUTOFFS

NOUS-HOSTED / 50 ATTEMPTED PER GROUP

GLM 5.3 Flash

44/50

26/50

Qwen3.8 Flash

31/50

24/50

LOGIN-BASED PROVIDER / FOUR SEPARATE CUTOFF GROUPS

MODEL	ROUTE	ATTEMPTED	GRADED	INFRA	NOT RUN
GPT-6 Luna	Direct	3 / 50	0	3	47
GPT-6 Luna	BloodHound MCP	3 / 50	0	3	47
GPT-5.6 Luna	Direct	3 / 50	0	3	47
GPT-5.6 Luna	BloodHound MCP	3 / 50	0	3	47

Each cutoff group: 3/50 attempted, 0 graded, 3 infrastructure errors, 47 unexecuted [\[4\]](#).
Not a zero reasoning score or a proven authentication failure. Direct/MCP scoring contracts differ.





The model-MCP matchup

FULL SCOREBOARD / CORRECT PER 50 SCHEDULED

MODEL	DIRECT	BH MCP	ARMADIN	STEVEN
GLM-5.3-Flash	39/50 	28/50 	16/50* 	31/50*
GPT-6 Luna	46/50 	49/50 	30/50* 	48/50*
GPT-5.6 Luna	42/50 	18/50† 	35/50* 	44/50*
DeepSeek-v4.1-Flash	29/50 	38/50* 	23/50* 	36/50*

† 23 tasks attempted; 27 unexecuted. * partial group

Seed 67 [\[1\]](#), [\[2\]](#), [\[3\]](#). One pass; all campaigns failed/partial. Shared questions ≠ identical settings or scoring contracts.





Where the tokens went

RESOURCE METERS / INPUT + OUTPUT TOKEN SHARE

Mint = input / Cream = output; each meter = 100%

TOTAL TOKENS

GLM 5.3 Flash



4.90M

95.1% input | 423 tool requests | 4 failed responses

Qwen3.8 Flash



6.24M

95.4% input | 504 tool requests | 10 failed responses



Sept 28 hosted BloodHound MCP groups; 50 attempted questions each [\[4\]](#). Mint=input / cream=output. No verified dollar costs.



Three passes, same questions

CORRECT / 150 SCHEDULED – 50 QUESTIONS × THREE PASSES

MODEL	DIRECT	BH MCP	ARMADIN	STEVEN
GPT-6 Luna	128/150 40 43 45 	145/150* 47 49 49 	95/150* 32 32 31 	117/150† 50 48 19
GLM-5.3-Flash	113/150 38 39 36 	73/150 23 26 24 	51/150* 15 19 17 	87/150* 27 30 30
DeepSeek-v4.1-Flash	99/150 32 35 32 	82/150 28 27 27 	73/150* 25 26 22 	104/150* 35 35 34

Small figures: correct per 50-task pass. * failed/partial rows. † 123/150 attempted.

Original interrupted. Local Qwen direct: 12/150 correct, 71 attempted; its MCP routes, Muse and Gemma not run [5].
 Shared questions do not make Direct and MCP scoring contracts identical.



Recovery is a separate result

SEPARATE RECOVERY CAMPAIGN / NOT REPLACEMENT ROWS

QWEN3.8 FLASH / CLOUD / THREE PASSES

ROUTE	CORRECT	ATTEMPTED
Direct	96/150	150/150
BH MCP	82/150	150/150
Armadin	26/150†	101/150
Steven	70/150*	150/150

GPT-6 LUNA / STEVEN

49/50*

One recovery repetition
50 attempted; cleanup failed

* failed/partial cleanup state. † Armadin left 49 evaluations unexecuted.

Recovery failed/partial [6]. Cloud Qwen differs from the original local model; GPT settings fingerprint changed.
No recovery row is silently spliced into the original campaign.





An API rejection isn't a wrong answer

HTTP 400 / REQUEST SIZE IS EVIDENCE, NOT A DIAGNOSIS

ORIGINAL / THREE PASSES

18

latest HTTP 400 outcomes
16 Armadin · 2 Steven

RECOVERY / QWEN × ARMADIN

16

latest HTTP 400 outcomes
49/150 evaluations not run

analyze_group_permissions

3,710,830 characters

Saved tool output · four matching failed attempts

MESSAGES + 97 TOOL SCHEMAS

≈ 4.49 MB

Reconstructed core · not full-request or wire bytes



No rejection body or size limit saved. Oversized-request cause remains unconfirmed.



Above: latest outcomes [7]. Recovery historical HTTP 400 attempts: 34 = 26 ultimately INFRA_ERROR + 8 COMPLETED [9].
Attempt counts are not distinct failed-question counts. No causal threshold or verified recovery fix.





The next questions for ORI

NEXT / AZURE

Azure BloodHound

Generate Azure BloodHound data.

Evaluate reasoning over cloud identity attack paths.

THEN / OPENGGRAPH

More graph sources

Add OpenGraph data from major collectors.

Start with GitHound and JamfHound.

EXPLORE / JUDGE

Model as a judge

Research Jev or an open-source alternative.

Check judgments against known answers first.

Broader model testing planned: Fable, Astra, Kimi (budget permitting).
Future goals and research — not completed capabilities.





The score is a starting point

Which model and MCP is best
at BloodHound?

Test the model + interface.

Read the mistakes and the cost.


Use the result to choose and improve.





Your next benchmark

<https://github.com/Specter0ps/offensive-reasoning-index>





Acknowledgements

Armadin team

[BloodHound MCP](#) | [Brett's blog](#)

Steven

[His BloodHound MCP](#)

MorDavid

[His MCP](#)

SpecterOps

[BloodHound](#) and help testing ORI internally

[Read: Chatting with Your Attack Paths](#)

Kyle / Outflank

Contributions to [my BloodHound MCP](#)

Thank you.



References

[1] GLM-5.3-Flash

Sept 29–30 / seed 67
ref-01-glm53.csv

[3] DeepSeek-v4.1-Flash

Sept 29–30 / seed 67
ref-03-deepseek.csv

[5] Three-repetition campaign

Seed 67 / interrupted; not three seeds
ref-05-three-repetitions.csv

[7] Provider HTTP 400 failures

Derived latest-outcome counts
ref-07-http-errors.csv

[9] Retained-tool and request size

Derived measurements / reconstructed core
ref-09-request-size.json

[2] GPT-6 / GPT-5.6 Luna

Sept 29–30 / seed 67
ref-02-luna.csv

[4] September 28 campaign

Scores, usage and login-provider cutoffs
ref-04-hosted.csv

[6] Separate recovery campaign

Seed 67 / failed, partial
ref-06-recovery.csv

[8] Local Qwen / five routes

Exact aggregate / coverage and outcomes
ref-08-local-qwen.csv

Repository: references/data/ · Original aggregates and sanitized derived diagnostics; raw traces excluded.